

Apollo AgriVerse

Engineering an Intelligent Agricultural Ecosystem

AgriIntel Division
Explainable Crop and Field Intelligence Vertical Slice

Engineering Principle: “Agriculture should not merely be monitored—it should be understood, anticipated, and empowered through intelligence.”

Document Type	Mid-Term Internship Technical Report
Work Package	AgriIntel Division
Team	Vijayalaxmi K. Sundalam — Technical Integration Nandini N. Naral — Product, Data Quality & Mining Sunaina S. Gaikwad — ML Baseline & Model Evaluation Dakshini A. Neel — Data & ML-Model Frontend
Supervisor	Akash Shivdas Chatake
Organization	Chatake Innworks Pvt. Ltd.
Internship Period	June 2026 – August 2026
Report Date	August 2026

MID-TERM STATUS
This report records the engineering state reached at the mid-term checkpoint. Completed, in-progress, under-evaluation and planned components are intentionally distinguished so that the report does not convert a roadmap into a false completion claim.

Preliminary. Declaration and Acknowledgement

This mid-term report presents the technical work undertaken by the AgriIntel engineering team within Apollo AgriVerse during the stated internship period. The report distinguishes implemented engineering work, simulated states, derived results, proposed architecture and remaining work so that the evidence boundary of the project remains explicit.

Student / Team Representative	Supervisor
_____ AgriIntel Engineering Team	_____ Akash Shivdas Chatake

Acknowledgement

The team expresses its sincere appreciation to **Akash Shivdas Chatake** for the architectural and domain-level feedback that materially refined the project. The team also acknowledges the Apollo Continuum workspace and the datasets, tools and technical resources used during the engineering process.

Executive. Executive Summary

Apollo AgriVerse is being developed as an intelligent agricultural ecosystem in which agricultural conditions are represented as evolving states rather than isolated readings. Within this larger ecosystem, AgriIntel is a focused engineering work package aimed at an explainable crop and field-intelligence vertical slice. The engineering principle is to **build one reviewable decision loop**: agricultural evidence is collected and validated, represented as system state, interpreted through intelligence mechanisms, and ultimately translated into an explainable decision or recommendation.

The mid-term work has progressed from domain research and requirement refinement into data engineering, synthetic sensor generation, state-based simulation, digital-twin architecture, ML exploration and frontend prototyping. The current architecture separates the dynamic state representation from predictive intelligence. This separation is deliberate: a state model answers *what condition the simulated or observed system is currently in and how it changes over time*, while a machine-learning model answers *what outcome is likely given the available evidence*. Treating these as different responsibilities improves auditability, supports temporal reasoning and avoids forcing a predictive model to become the sole source of system memory.

The project is currently positioned between subsystem development and end-to-end integration. The remaining work is therefore not framed as a list of unrelated features; it is the completion and validation of the decision loop through backend integration, ML evaluation, state synchronization, frontend binding, evidence capture and final deployment.

MID-TERM ENGINEERING THESIS

The key achievement is not the number of individual modules created. It is the establishment of a coherent architecture in which data, state, simulation, intelligence and interface can eventually participate in one traceable agricultural decision loop.

Contents. Table of Contents

Preliminary. Declaration and Acknowledgement	2
Acknowledgement	2
Executive. Executive Summary	3
Contents. Table of Contents	4
1. Introduction	7
1.1 Project Background	7
1.2 Problem Statement	7
1.3 Project Vision	7
1.4 Objectives	7
2. Agricultural Problem and Domain Analysis	7
2.1 Why Agricultural Intelligence Is a Dynamic Problem	7
2.2 Soil–Crop–Environment Relationship	7
2.3 Crop Lifecycle Representation	8
3. Research and Requirement Engineering	8
3.1 Research Direction	8
3.2 Functional Requirements	8
3.3 Non-Functional Requirements	8
4. Project Evolution and Scope Refinement	9
4.1 Initial Technical Direction	9
4.2 Supervisor Feedback as a Requirements Refinement	9
4.3 Effect on Architecture	9
5. Apollo AgriVerse System Architecture	9
5.1 System Hierarchy and AgriIntel Boundary	9
5.2 The Digital Twin as a State Representation	10
5.3 Why State Models Were Chosen	11
5.4 Why Not Use Machine Learning for Everything?	11
5.5 Backend and Storage Architecture	12
6. Data Acquisition and Engineering	12
6.1 Data Strategy	12
6.2 Dataset Inventory	13
6.3 Synthetic Data Generation	13
6.4 Data Quality and Integration	13

7. State Simulation and Digital Twin Development	13
7.1 State Models: Definition	14
7.2 Soil State Model	14
7.3 Hydrogel State Model	14
7.4 Mulching State Model	14
7.5 Crop Lifecycle and GDD	14
7.6 Why State Simulation Is Valuable Before Real-Time IoT Data	14
8. Machine Learning and Intelligence	14
8.1 Role of Machine Learning	14
8.2 Why ML Is a Layer Rather Than the Whole System	15
8.3 Model Selection Philosophy	15
8.4 Current Model Status	15
9. User Interface and Visualization	15
9.1 Purpose of the Frontend	16
9.2 Current Dashboard Prototype	16
9.3 Frontend–State–ML Relationship	16
10. Technology Stack and Development Environment	16
11. Development Methodology and Team Workflow	17
11.1 Workstream Structure	17
11.2 Iterative Engineering Cycle	17
12. Mid-Term Progress and Current Status	17
13. Challenges, Engineering Decisions and Resolutions	18
14. Preliminary Results and Validation	19
14.1 What Has Been Validated	19
14.2 Evidence Classes	19
14.3 Current Validation Limitations	19
15. Remaining Work and Deployment Roadmap	19
15.1 Immediate Engineering Priorities	20
15.2 Deployment Readiness Criteria	20
16. Future Scope	20
17. Interim Conclusion	21
18. References	21

Appendices. Technical Evidence Appendices	21
Appendix A - Evidence and Claim Classification	21
Appendix B - Representative State Schema	22
Appendix C - Frontend Evidence	22

The contents list is generated from the actual chapter and subsection anchors in the PDF. Chapter entries are also written into the PDF outline/bookmark tree, allowing direct navigation.

1. Introduction

1.1 Project Background

Agricultural systems are inherently dynamic. Soil moisture changes with weather and irrigation; nutrient availability changes through uptake and intervention; crop development progresses through biological stages; and environmental stress can alter the trajectory of the crop before a conventional monitoring dashboard visibly reports a problem. A useful intelligence system therefore needs more than a collection of sensor values. It needs a representation of *state*, a mechanism for updating that state, and a reasoning layer capable of interpreting the state.

1.2 Problem Statement

The engineering problem addressed by AgriIntel is the construction of an explainable agricultural intelligence slice capable of connecting heterogeneous agricultural evidence to a continuously evolving representation of the farm and, ultimately, to a reviewable decision. The challenge is not simply prediction. It includes data quality, temporal alignment, state representation, simulation, model selection, integration and communication of the result to the user.

1.3 Project Vision

The long-term vision is a system in which a farmer or agronomist can move from a selected farm and crop context to a meaningful explanation of current conditions, expected development, emerging stress and appropriate intervention. The mid-term implementation intentionally treats this as an engineering path rather than claiming that the final product already exists.

1.4 Objectives

- Establish a coherent agricultural data foundation covering soil, weather, crop, farm and intervention context.
- Represent important agricultural variables as dynamic states that can evolve with time and new evidence.
- Develop a digital-twin-oriented architecture that separates state management from predictive intelligence.
- Evaluate where machine learning is appropriate and where explicit state-transition logic is more defensible.
- Create an explainable interface through which data and model outputs can eventually become a reviewable decision.
- Reach an integrated, testable vertical slice before final deployment rather than prematurely claiming a complete platform.

2. Agricultural Problem and Domain Analysis

2.1 Why Agricultural Intelligence Is a Dynamic Problem

Agricultural variables are coupled across time. Soil moisture affects root-zone conditions; weather changes evaporative demand; interventions such as hydrogel and mulching alter water retention and surface conditions; and crop stage changes the significance of the same environmental measurement. Consequently, a value such as 30% soil moisture cannot be interpreted in isolation. Its meaning depends on crop, soil, recent history, weather, intervention state and lifecycle stage.

2.2 Soil–Crop–Environment Relationship

The project's refined direction therefore treats soil, crop and environment as a coupled system. The soil layer provides a physical and chemical context; weather provides external forcing; the crop lifecycle determines biological demand; and intervention layers modify the state trajectory. This relationship is the reason a simple static dashboard is insufficient for the intended intelligence layer.

2.3 Crop Lifecycle Representation

Calendar days are a weak representation of biological progress because two seasons with different thermal conditions can produce different developmental rates. Growing Degree Days provide a more meaningful intermediate representation by accumulating temperature exposure relative to a crop base temperature. In the current architecture, GDD is therefore treated as a state-transition driver for crop lifecycle simulation rather than merely another dashboard metric.




Figure 2.1 — Illustrative crop lifecycle representation using cumulative Growing Degree Days.

3. Research and Requirement Engineering

3.1 Research Direction

The research phase examined agricultural variables, digital-twin concepts, state modelling, machine learning, data engineering, crop lifecycle representation and interface requirements. The purpose of this research was not to select the most sophisticated technology available; it was to determine which technical abstraction best matched the behaviour that the system must represent.

3.2 Functional Requirements

Requirement	Engineering interpretation	Current status
Agricultural context	Farm, soil, weather and crop context must be represented together.	COMPLETED / IN PROGRESS
Temporal representation	System state must evolve as evidence changes.	IN PROGRESS
Crop lifecycle	Growth stage should be represented continuously rather than only as a static label.	COMPLETED
Intelligence	Predictive models should consume engineered state/evidence where justified.	UNDER EVALUATION
Explainability	User should be able to understand the evidence behind a result.	IN PROGRESS
Integration	Frontend, backend, state engine and ML must form one demonstrable loop.	PLANNED

3.3 Non-Functional Requirements

- Traceability: system outputs should be attributable to identifiable input data and state.
- Maintainability: simulation, data processing, ML and UI responsibilities should remain modular.
- Reproducibility: generated states and model experiments should be repeatable from controlled inputs.

- Extensibility: new crops, soil types, interventions and models should be addable without redesigning the entire system.
- Evidence discipline: simulated, observed, predicted and planned states must not be presented as equivalent evidence.

4. Project Evolution and Scope Refinement

4.1 Initial Technical Direction

The early project direction was primarily technology-led. The team had begun exploring data pipelines, models, simulation and interface concepts, but the agricultural decision context was not yet sufficiently explicit. This created a risk: individual components could be technically functional while the overall system remained weak as an agricultural solution.

4.2 Supervisor Feedback as a Requirements Refinement

Project review introduced a stronger domain perspective. The team was pushed to ask what soil the farm has, which crop or variety is appropriate, what that soil requires, how environmental conditions affect the crop, how interventions alter field conditions and how the crop should be followed from establishment through maturity. This changed the engineering question from *how do we display or predict a variable?* to *how do we represent and reason about an evolving agricultural system?*

- Soil type and regional variation became first-class inputs.
- Crop and variety selection became part of the agricultural context.
- Crop lifecycle became a dynamic component rather than a static attribute.
- Intelligent soil, hydrogel and mulching concepts were connected to state changes.
- India-wide weather and soil data were considered to improve geographic generality.
- The digital twin was reframed as a state-and-simulation architecture rather than a visual 3D object.

4.3 Effect on Architecture

The consequence was an architectural shift toward explicit state representation, synchronization and separation of concerns. The revised system can therefore accommodate evidence collection, state evolution, predictive intelligence and user interaction without forcing every concern into a single model.

Engineering interpretation

The scope change is treated as a requirements-maturation event. The important result is not that the original direction changed; it is that the team identified a gap before final deployment and restructured the architecture while the project was still sufficiently flexible to absorb the change.

5. Apollo AgriVerse System Architecture

5.1 System Hierarchy and AgriIntel Boundary

AgriIntel is positioned as a focused work package within Apollo AgriVerse. Its engineering target is an explainable crop and field-intelligence vertical slice, with the working principle of building one reviewable decision loop. The architecture therefore emphasizes the path from agricultural evidence to a state representation, intelligence and an explainable decision.





Figure 5.1 — Apollo AgriVerse system hierarchy, AgriIntel boundary and decision-loop architecture.

5.2 The Digital Twin as a State Representation

In this project, the digital twin is not defined as a 3D visualization. A digital twin becomes useful when a computational representation maintains a meaningful correspondence with a physical or simulated system over time. For AgriIntel, that correspondence is expressed through variables such as soil moisture, nutrient condition, root-zone thermal state, hydrogel capacity, mulch degradation and crop lifecycle stage.

The twin therefore has three responsibilities: **representation** of the current condition, **transition** from the current condition to the next state, and **synchronization** when new evidence becomes available. The intelligence layer can then consume a consistent state instead of reconstructing temporal context independently every time it makes a prediction.



The state model is an explicit memory of the system condition; it is not a prediction algorithm.

Figure 5.2 — State-model abstraction used to represent evolving agricultural conditions.

5.3 Why State Models Were Chosen

A state model was selected because several project variables are not simply labels or prediction targets; they represent conditions that persist and evolve. Soil moisture has memory. Crop lifecycle has progression. Mulch degrades. Hydrogel capacity changes as water is absorbed and released. A state representation makes these temporal dependencies explicit.

Alternative	Advantage	Limitation for current problem	Decision
Pure static dataset	Simple and easy to query.	Cannot represent evolving system condition.	Rejected as primary representation
Pure ML model	Can learn nonlinear relationships.	Does not automatically provide an explicit, auditable system state.	Used as intelligence layer, not sole state
Full mechanistic model	Strong physical interpretation.	Requires detailed calibration, parameters and observability that are not yet available.	Future reference / selective rules
State-transition model	Explicit temporal memory and transparent transitions.	Requires carefully designed transition logic.	Selected for current twin foundation
Hybrid state + ML	Combines explicit state with predictive intelligence.	More integration effort.	Target architecture

5.4 Why Not Use Machine Learning for Everything?

Using ML to predict every intermediate state would create an unnecessary dependence on labelled historical trajectories. The project does not currently have sufficiently rich, continuously observed field histories for every state variable. In addition, an ML model can learn a statistical relationship without exposing the internal state transition that produced it. That is undesirable when the system must explain why a crop is considered stressed or why an intervention changed the simulated field condition.

This does not make state models a replacement for ML. Instead, the two solve different problems. The state model maintains the system's evolving context; ML can learn relationships that are difficult to express explicitly, such as yield risk, stress classification or intervention demand. The resulting architecture is therefore hybrid by design.





Figure 5.3 — Complementary roles of explicit state models and machine-learning models.

5.5 Backend and Storage Architecture

The backend is the integration boundary between the user interface and the computational services. Its role is not to duplicate the digital twin or embed all modelling logic inside HTTP routes. Instead, the backend should expose controlled operations such as reading the current twin state, updating state from validated input, requesting a model inference and returning an explanation-ready response.

SQLite is currently useful as a lightweight state/history store because it is easy to inspect, version alongside the prototype and query from Python services without introducing infrastructure overhead prematurely. It is not being treated as the final production database architecture. The current choice is appropriate for a reviewable prototype where state persistence and reproducibility matter more than distributed scale.



The backend exposes controlled operations, not a duplicate of the twin or the entire UI.

Figure 5.4 — Backend integration boundary between frontend, state services, ML inference and stored artifacts.

6. Data Acquisition and Engineering

6.1 Data Strategy

The data foundation combines collected/real-world agricultural datasets with synthetic data generated for system simulation. The distinction is important. Real datasets provide evidence about observed relationships; synthetic datasets provide controlled coverage for testing state transitions, integration logic and edge cases that are difficult to obtain from limited field observations.




Figure 6.1 — Multi-source data processing and integration pipeline.

6.2 Dataset Inventory

Data component	Purpose	Representative variables	Current role
Soil	Agronomic context	pH, N, P, K, texture, organic indicators	State initialization / intelligence
Weather	Environmental forcing	Temperature, humidity, rainfall and related climate variables	State transition / lifecycle
Crop	Crop suitability and context	Crop identity, requirements and mappings	Context / recommendation
Farm metadata	Farm identity and configuration	Farm ID, location/context, configuration	Twin identity
Sensor stream	Dynamic telemetry representation	Moisture, temperature, humidity and other sensor-like signals	State synchronization
Hydrogel	Intervention behaviour	Retention/release-related variables	State transition
Mulching	Surface intervention	Degradation/cooling/evaporation-related variables	State transition
Crop lifecycle	Phenological progression	GDD, growth stage and stage transitions	Lifecycle state

6.3 Synthetic Data Generation

Synthetic generation is used where the project needs a controlled stream of state changes. The purpose is not to pretend that synthetic observations are equivalent to field measurements. Instead, a generator allows the team to test whether the twin engine behaves correctly when moisture declines, environmental conditions change, interventions modify the trajectory or crop lifecycle thresholds are crossed.

6.4 Data Quality and Integration

Before integration, datasets require schema alignment, type normalization, missing-value handling, unit checks, outlier review and temporal/geographical consistency checks. The master representation should preserve source semantics rather than flattening every variable into an undifferentiated table. This is particularly important when future image data is introduced.




Figure 6.2 — Recommended separation of tabular, state and image modalities while preserving semantic linkage.

7. State Simulation and Digital Twin Development

7.1 State Models: Definition

A state model describes the condition of a system at a particular time using a set of variables that are sufficient for the intended simulation and decision task. Conceptually, the system can be expressed as $S_t = \{\text{soil, environment, intervention, crop}\}$. A transition function then maps the current state and new evidence to a next state: $S_{t+1} = F(S_t, E_{t+1}, A_t)$, where E is new evidence and A represents intervention effects.

The advantage is architectural clarity. The system can ask what changed between two timestamps, which variables caused the transition, and which state should be exposed to downstream intelligence. This is substantially different from storing only the latest sensor row.

7.2 Soil State Model

The soil state represents a dynamic medium rather than a fixed metadata record. Relevant state variables may include moisture, nutrient condition and thermal stress. Soil type and baseline chemistry initialize the state, while weather, crop demand and interventions influence subsequent transitions.

7.3 Hydrogel State Model

The hydrogel component is represented as an intervention whose capacity changes as water is retained and released. Its purpose in the twin is therefore not simply to store a material property; it provides a mechanism through which an intervention can alter the simulated moisture trajectory.

7.4 Mulching State Model

Mulching is represented as a time-dependent intervention because its effect is expected to change as the material degrades. The state therefore captures the remaining intervention effectiveness and its influence on surface conditions such as evaporation or temperature suppression.

7.5 Crop Lifecycle and GDD

The crop lifecycle state uses accumulated heat exposure as a transition signal. A simplified conceptual relation is $GDD = \max(((T_{max} + T_{min}) / 2) - T_{base}, 0)$. Cumulative GDD can then be compared with crop-specific stage thresholds. The important architectural point is that the crop stage becomes a state variable that influences the interpretation of other variables.

7.6 Why State Simulation Is Valuable Before Real-Time IoT Data

A real-time field twin cannot be validated merely by connecting sensors. The software must first demonstrate that it can accept a sequence of observations, update its state deterministically, retain history and expose the resulting state to downstream components. Synthetic simulation provides this controlled environment. It is therefore a development and validation mechanism, not a substitute for field evidence.

Alternative architectures considered

Rule-only: highly transparent but limited for nonlinear learned relationships. **ML-only:** predictive but weak as an explicit state memory. **Full mechanistic:** scientifically rich but calibration-heavy. **Hybrid state + ML:** explicit state provides temporal context while ML handles learned relationships. The current project follows the fourth approach.

8. Machine Learning and Intelligence

8.1 Role of Machine Learning

Machine learning is introduced where the relationship between agricultural evidence and the desired output is difficult to encode through deterministic transition rules alone. Candidate tasks include yield estimation, stress risk, intervention demand and other derived intelligence. The model should consume well-defined features and return a result that can be placed back into the decision loop.

8.2 Why ML Is a Layer Rather Than the Whole System

The distinction is important for engineering reliability. A model can predict an outcome without maintaining a canonical representation of the farm's current condition. If the UI, simulator and ML pipeline each reconstruct state differently, the same farm can produce inconsistent interpretations. The twin therefore provides a common state context, while ML provides learned inference.

8.3 Model Selection Philosophy

The project should not select a model because it is more complex. Tabular agricultural data often supports strong tree-based baselines, while sequence models become relevant when the available data contains sufficiently rich and regular temporal histories. A model should therefore be promoted only after the data volume, temporal structure, validation protocol and target definition justify it.

Candidate approach	Why it is useful	Why it is not automatically sufficient
Linear / simple baseline	Transparent reference point; identifies whether complexity is necessary.	May miss nonlinear interactions.
Tree-based model	Strong for heterogeneous tabular features and nonlinear interactions.	Does not inherently maintain a physical state or chronology.
Sequence model	Can learn temporal dependencies when adequate sequences exist.	Requires sufficient, correctly sampled temporal evidence; can overfit.
State + ML hybrid	Maintains explicit state and adds learned prediction where useful.	Requires more careful integration and validation.

8.4 Current Model Status

Two concrete ML pipelines have been implemented and evaluated at the current checkpoint. The first is a grape-yield regression pipeline using an 800-row, 9-column agronomy dataset. Random Forest and Gradient Boosting regressors were compared using 5-fold cross-validation; Random Forest was retained as the better of the two candidates, but its held-out R^2 of -0.0942 does not support treating the yield model as production-ready. The second pipeline uses a 1,000-row, 24-column digital-twin telemetry dataset to predict hydrogel water-storage percentage. Its Gradient Boosting model achieved a test RMSE of 0.9484%, MAE of 0.1456% and R^2 of 0.9900. These results are reported as notebook-level evaluation evidence; they should not be interpreted as field validation.

Model / pipeline	Target	Data / split	Current evaluation	Mid-term interpretation
Grape Yield Prediction — Random Forest Regressor	yield_tons_ha	800 rows; 640 train / 160 test; 5-fold CV	CV R^2 = -0.0944 ± 0.0299; Test RMSE = 0.6134 tons/ha; MAE = 0.4569 tons/ha; R^2 = -0.0942	Baseline retained for comparison, not production-ready.
Grape Telemetry & Hydrogel Trigger — Gradient Boosting Regressor	hydrogel_water_storage_pct	1,000 rows; 800 train / 200 test	Test RMSE = 0.9484%; MAE = 0.1456%; R^2 = 0.9900	Strong notebook-level result; requires broader validation before deployment.


The important architectural conclusion is that ML now occupies the predictive-intelligence layer, while the state models remain responsible for explicit temporal representation and transition. The two mechanisms are therefore complementary rather than competing replacements. The next step is to validate the models against the intended data-generation assumptions and then connect verified artifacts to the backend inference layer.

9. User Interface and Visualization


9.1 Purpose of the Frontend

The frontend is not merely a visualization layer. It is the communication surface through which the state of the agricultural system, model outputs and recommended actions become understandable to the intended user. A dashboard that displays many charts without explaining their relationship to a decision would reproduce the same fragmentation the project is attempting to address.

9.2 Current Dashboard Prototype




Dashboard Panel 1




Dashboard Panel 3




Dashboard Panel 5



Dashboard Panel 7



Dashboard Panel 9



Dashboard Panel 12

Selected frontend panels from the current prototype; the complete 12-panel interface is retained as project evidence but is not reproduced in full here.

The screenshots above are real prototype panels already produced by the team. They are intentionally presented as evidence of interface development, not as proof of final backend integration. The complete panel set contains 12 detected screenshots in the asset directory; only a representative subset is included to maintain information density.

9.3 Frontend-State-ML Relationship

A future integrated dashboard should not independently calculate business logic in each panel. The intended relationship is: the backend retrieves or updates the canonical twin state; the intelligence layer produces model outputs from controlled features; the frontend renders those outputs with their context. This reduces the risk that two dashboard panels display contradictory interpretations of the same farm state.

10. Technology Stack and Development Environment

Layer	Technology / Tool	Purpose
Core language	Python	Data processing, simulation, orchestration and ML.
Data engineering	Pandas / NumPy	Cleaning, transformation, feature preparation and analysis.
ML	Scikit-learn / verified project models	Baseline modelling and evaluation where applicable.
Simulation	Python state engines	Temporal state transitions and synthetic streams.
Storage	SQLite	Prototype state/history persistence and inspectable local storage.
Frontend	Streamlit / React Prototype	Dashboard and user-facing visualization.
Backend	SQLite / Relational schema	API/integration boundary.
Version control	Git / GitHub	Source control and collaborative engineering.
Reporting	ReportLab / Matplotlib	Automated technical report and evidence figures.

Technology selection at the mid-term stage is intentionally pragmatic. Lightweight tools reduce infrastructure overhead while the architecture is still changing. Production technologies should be finalized only after the decision loop and validation requirements are stable.

11. Development Methodology and Team Workflow

11.1 Workstream Structure

Workstream	Responsibility
Product, Data Quality & Mining	Field-record quality, mining workflow, scenarios and acceptance criteria.
Technical Integration	Repository health, environment, integration and shared pipeline work.
ML Baseline & Model Evaluation	Baseline modelling, evaluation and model-selection evidence.
Data & ML-Model Frontend	Data/model interface, dashboard and component demonstration.

11.2 Iterative Engineering Cycle

The project follows an iterative cycle: research and requirement clarification → data preparation → subsystem implementation → evidence generation → review → architectural refinement. This is particularly important because the supervisor feedback changed the project's agricultural framing while implementation was already underway.

- Keep workstreams modular but expose clear interfaces between them.
- Use Git/GitHub to retain implementation history and integration points.
- Generate evidence during development rather than reconstructing it at the end.
- Treat model evaluation and state validation as separate engineering activities.
- Record limitations explicitly when evidence is synthetic or incomplete.

12. Mid-Term Progress and Current Status




Figure 12.1 — Mid-term engineering maturity view across the principal work areas. Percentages are project-planning maturity bands, not performance metrics.

Subsystem	Status	Mid-term achievement	Next step
Research & scope	COMPLETED	Agricultural problem and AgriIntel boundary clarified.	Maintain evidence boundary.
Data engineering	IN PROGRESS	Soil, weather, crop, farm and synthetic components prepared.	Complete integrated master representation.
State simulation	COMPLETED / IN PROGRESS	Core soil, hydrogel, mulch and lifecycle state simulations explored.	Unify state transitions and synchronization.
Digital Twin engine	IN PROGRESS	State-oriented architecture and synchronizer concept established.	Connect validated data streams and persistence.
ML intelligence	UNDER EVALUATION	Model pipeline exploration and baseline work underway.	Verify metrics and select defensible models.
Frontend	IN PROGRESS	Real dashboard panels prototyped.	Bind panels to backend/state/ML outputs.
Decision loop	PLANNED	Interfaces between components identified.	Demonstrate one complete reviewable loop.
Deployment	PLANNED	Deployment requirements identified.	Validation, hardening and release.

13. Challenges, Engineering Decisions and Resolutions

Challenge	Why it mattered	Engineering response	Status
Project initially too technology-centric	Risk of building components without a strong agricultural decision context.	Refined scope around soil–crop–environment intelligence and one reviewable decision loop.	RESOLVED
Limited continuous field evidence	Pure data-driven state modelling would require more labelled temporal data.	Use explicit state models and controlled synthetic streams while treating ML as a separate intelligence layer.	IN PROGRESS
Multiple datasets with different semantics	Naive merging could destroy source meaning and temporal relationships.	Normalize schemas and retain semantic boundaries before integration.	IN PROGRESS
Digital twin terminology can be overstated	A visualization alone is not a continuously synchronized twin.	Define twin around state, transition, synchronization and evidence.	RESOLVED

Challenge	Why it mattered	Engineering response	Status
Frontend development can become disconnected from backend logic	Panels may display inconsistent or duplicated calculations.	Use backend/state services as the integration boundary.	IN PROGRESS
Roadmap vs evidence	Future features can easily be mistaken for completed work.	Use explicit status labels and evidence-bound claims.	RESOLVED

14. Preliminary Results and Validation

14.1 What Has Been Validated

At the mid-term stage, validation is primarily architectural and subsystem-level. The team has demonstrated that the required data components can be represented, that synthetic streams can be generated for controlled testing, that state transitions can be simulated for the major conceptual pillars, and that a real dashboard prototype exists. These are meaningful engineering milestones, but they do not yet constitute final field validation.

14.2 Evidence Classes


Evidence class	Meaning in this project	How it should be reported
Observed	Values originating from actual collected/public datasets.	Use source and dataset identity.
Derived	Values calculated from observed data or rules.	Show transformation or formula.
Simulated	Values generated by state/synthetic models.	Label explicitly as simulated.
Predicted	Output of an ML model.	Report model, input context and evaluation protocol.
Proposed	Architecture or future functionality not yet implemented.	Keep in roadmap / future scope.

14.3 Current Validation Limitations

- Continuous physical sensor synchronization has not yet been established as a production field deployment.
- Synthetic streams cannot establish real-world model generalization.
- Final ML metrics must be taken from retained training/evaluation artifacts.
- State-transition parameters require calibration against real agricultural observations.
- End-to-end backend–twin–ML–frontend validation remains a remaining integration milestone.

15. Remaining Work and Deployment Roadmap

The remaining work is organized around completing the reviewable decision loop rather than simply adding more features. First, the state and data layers must be integrated and made reproducible. Second, verified ML models must be connected to those states. Third, backend services must expose controlled operations to the frontend. Finally, the complete loop must be tested under representative scenarios before deployment.



Roadmap is an evidence path, not a claim that future stages are already complete.

Figure 15.1 — Mid-term to final deployment roadmap. The figure intentionally contains no embedded figure number; numbering is generated by the report.

15.1 Immediate Engineering Priorities

- Finalize the integrated master data/state representation.
- Complete twin state synchronization and persistence.
- Verify actual ML models, metrics and evaluation protocol.
- Connect backend endpoints/services to state and inference layers.
- Bind representative frontend panels to live system outputs.
- Demonstrate one complete reviewable decision loop.

15.2 Deployment Readiness Criteria

Criterion	Required evidence
Data integrity	Validated schemas, source semantics and reproducible preprocessing.
State integrity	Deterministic state transitions and persisted history.
Model integrity	Verified model artifact, evaluation protocol and retained metrics.
API integrity	Controlled input/output contracts and error handling.
UI integrity	Panels display canonical state/model outputs without duplicated logic.
Decision integrity	Recommendation can be traced to evidence, state and model context.
Operational integrity	Deployment environment, logging and failure behaviour verified.

16. Future Scope

The architecture leaves room for several future extensions once the current decision loop is validated. These include real IoT telemetry, image-based crop and canopy diagnostics, richer crop-variety intelligence, more advanced predictive models, uncertainty estimation, continuous learning, broader regional coverage and production-grade database/API infrastructure.

Image intelligence is intentionally treated as a separate modality rather than being forced into the current tabular master dataset. A future multimodal architecture can associate an image with farm ID, timestamp, crop stage and twin state while preserving the image itself as an image artifact. This avoids a technically incorrect flattening of image data into ordinary numeric columns.

17. Interim Conclusion

At the mid-term checkpoint, Apollo AgriVerse — AgriIntel has moved beyond initial concept exploration into a structured engineering architecture. The team has established the agricultural problem context, refined the scope through review, prepared heterogeneous data components, developed state-based simulations for the principal agricultural pillars, explored the digital-twin state/synchronization architecture and produced a real frontend prototype.

The central architectural decision is to separate **state representation** from **predictive intelligence**. This provides the system with an explicit temporal memory while allowing machine learning to be used where learned relationships provide value. The next stage is therefore integration and evidence: one reviewable decision loop must be executed end-to-end, validated, explained and hardened before deployment.

Mid-term position

The project is not presented as a finished agricultural product. It is presented as an engineering system whose core abstractions, data foundation, simulation mechanisms and interface direction are sufficiently established to enter the integration and validation phase.

18. References

- [1] Apollo Continuum, "Apollo AgriVerse — Mission II: AgriIntel Work Package," official project workspace.
- [2] Apollo AgriVerse project working documentation, "Digital Twin Architecture & ML Manual," internal project document.
- [3] Internal AgriIntel datasets, notebooks, simulation scripts and frontend artifacts, accessed during the internship.
- [4] MindForgeAI Project Centralization Blueprint V1
- [5] Apollo AGriVerse Strategic System Blueprint.

Appendices. Technical Evidence Appendices

Appendix A - Evidence and Claim Classification

This appendix provides a compact evidence taxonomy intended to prevent overstatement in the final report and deployment review.

Claim type	Permitted evidence	Example
Implemented	Executable code / working prototype / retained artifact.	State simulator executes and persists output.
Validated	Repeatable test or evaluation with documented protocol.	Model metric on a defined test split.
Simulated	Synthetic or rule-generated evidence.	Hydrogel state trajectory from controlled generator.
Prototype	UI or subsystem demonstration not yet integrated end-to-end.	Current dashboard panel.
Planned	Design/roadmap only.	Production deployment or future IoT integration.

Appendix B - Representative State Schema

State group	Representative fields	Purpose
Farm identity	farm_id, region/context	Identifies the simulated/observed farm.
Soil	pH, N, P, K, moisture	Initial and evolving substrate condition.
Environment	temperature, humidity, rainfall	External forcing variables.
Intervention	hydrogel capacity, mulch degradation	Tracks intervention state.
Crop	crop_id, variety, growth stage, cumulative GDD	Biological lifecycle context.
Time	timestamp / simulation step	Maintains temporal ordering.
Intelligence	prediction, risk, recommendation context	Stores derived decision evidence where required.

Appendix C - Frontend Evidence

The current project asset directory contains real dashboard screenshots. Only a representative subset is included in Chapter 9 to maintain report density. The full set remains part of the project evidence package.

Panel	Evidence status
Dashboard Panel 1	Available in assets/images
Dashboard Panel 2	Available in assets/images
Dashboard Panel 3	Available in assets/images
Dashboard Panel 4	Available in assets/images
Dashboard Panel 5	Available in assets/images
Dashboard Panel 6	Available in assets/images
Dashboard Panel 7	Available in assets/images
Dashboard Panel 8	Available in assets/images
Dashboard Panel 9	Available in assets/images
Dashboard Panel 10	Available in assets/images
Dashboard Panel 11	Available in assets/images
Dashboard Panel 12	Available in assets/images